

EECS201000

Introduction to Programming Laboratory

Homework 4: Blocked All-Pairs Shortest Path

Due: Aug 3, 2017, 8AM

1 GOAL

This assignment helps you get familiar with CUDA on multi-GPU environment by implementing a blocked all-pairs shortest path algorithm. We encourage you to optimize your program by exploring different optimizing strategies for optimization points.

2 PROBLEM DESCRIPTION

In this assignment, you are asked to modify sequential Floyd-Warshall algorithm to a parallelized CUDA version which take advantages of multiple GPUs.

Given an $N \times N$ matrix $W = [w(i, j)]$ where $w(i, j) \geq 0$ represents the distance (weight of the edge) from a vertex i to a vertex j in a **simple directed graph** with N vertices. We define an $N \times N$ matrix $D = [d(i, j)]$ where $d(i, j)$ denotes the shortest-path distance from a vertex i to a vertex j . Let $D^{(k)} = [d^{(k)}(i, j)]$ be the result which all the intermediate vertices are in the set $\{1, 2, \dots, k\}$.

We define $d^{(k)}(i, j)$ as follows:

$$d^{(k)}(i, j) = \begin{cases} w(i, j) & \text{if } k = 0; \\ \min \left(d^{(k-1)}(i, j), d^{(k-1)}(i, k) + d^{(k-1)}(k, j) \right) & \text{if } k \geq 1. \end{cases}$$

The matrix $D^{(N)} = d^{(N)}(i, j)$ gives the answer to the APSP problem.

In the blocked APSP algorithm, we partition D into $[N/B] \times [N/B]$ blocks of $B \times B$ submatrices. The number B is called **blocking factor**. For instance, we divide a 6×6 matrix into 3×3 submatrices (or blocks) by $B = 2$.




Figure 1: Divide a matrix by $B = 2$

Blocked version of Floyd-Warshall algorithm will perform $\lceil N/B \rceil$ rounds, and each round is divided into 3 phases. It performs B iterations in each phase.

Assumes a block is identified by its index (I, J) , where $1 \leq I, J \leq \lceil N/B \rceil$. The block with index (I, J) is denoted by $D_{(I,J)}^{(k)}$.

In the following explanation, we assume $N = 6$ and $B = 2$. The execution flow is described step by step as follows:

- **Phase 1: Self-dependent blocks**

In the K -th iteration, the 1st phase is to compute $B \times B$ pivot block $D_{(K,K)}^{(K \times B)}$.

For instance, in the 1st iteration, $D_{1,1}^{(2)}$ is computed as follows:

$$\begin{aligned}
 d^{(1)}(1,1) &= \min \left(d^{(0)}(1,1), d^{(0)}(1,1) + d^{(0)}(1,1) \right) \\
 d^{(1)}(1,2) &= \min \left(d^{(0)}(1,2), d^{(0)}(1,1) + d^{(0)}(1,2) \right) \\
 d^{(1)}(2,1) &= \min \left(d^{(0)}(2,1), d^{(0)}(2,1) + d^{(0)}(1,1) \right) \\
 d^{(1)}(2,2) &= \min \left(d^{(0)}(2,2), d^{(0)}(2,1) + d^{(0)}(1,2) \right) \\
 d^{(2)}(1,1) &= \min \left(d^{(1)}(1,1), d^{(1)}(1,2) + d^{(1)}(2,1) \right) \\
 d^{(2)}(1,2) &= \min \left(d^{(1)}(1,2), d^{(1)}(1,2) + d^{(1)}(2,2) \right) \\
 d^{(2)}(2,1) &= \min \left(d^{(1)}(2,1), d^{(1)}(2,2) + d^{(1)}(2,1) \right) \\
 d^{(2)}(2,2) &= \min \left(d^{(1)}(2,2), d^{(1)}(2,2) + d^{(1)}(2,2) \right)
 \end{aligned}$$

Note that result of $d^{(2)}$ depends on the result of $d^{(1)}$ and therefore cannot be computed in parallel with the computation of $d^{(1)}$.

- **Phase 2: Pivot-row and pivot-column blocks**

In the K -th iteration, it computes all $D_{(h,K)}^{(K \times B)}$ and $D_{(K,h)}^{(K \times B)}$ where $h \neq K$.

The result of pivot-row/pivot-column blocks depend on the result in Phase 1 and itself

For instance, in the 1st iteration, the result of $D_{(1,3)}^{(2)}$ depends on $D_{(1,1)}^{(2)}$ and $D_{(1,3)}^{(0)}$:

$$d^{(1)}(1,5) = \min \left(d^{(0)}(1,5), d^{(2)}(1,1) + d^{(0)}(1,5) \right)$$

$$d^{(1)}(1,6) = \min \left(d^{(0)}(1,6), d^{(2)}(1,1) + d^{(0)}(1,6) \right)$$

$$d^{(1)}(2,5) = \min \left(d^{(0)}(2,5), d^{(2)}(2,1) + d^{(0)}(1,5) \right)$$

$$d^{(1)}(2,6) = \min \left(d^{(0)}(2,6), d^{(2)}(2,1) + d^{(0)}(1,6) \right)$$

$$d^{(2)}(1,5) = \min \left(d^{(1)}(1,5), d^{(2)}(1,2) + d^{(1)}(2,5) \right)$$

$$d^{(2)}(1,6) = \min \left(d^{(1)}(1,6), d^{(2)}(1,2) + d^{(1)}(2,6) \right)$$

$$d^{(2)}(2,5) = \min \left(d^{(1)}(2,5), d^{(2)}(2,2) + d^{(1)}(2,5) \right)$$

$$d^{(2)}(2,6) = \min \left(d^{(1)}(2,6), d^{(2)}(2,2) + d^{(1)}(2,6) \right)$$

- **Phase 3: Other blocks**

In the K -th iteration, it computes all $D_{(h_1,h_2)}^{(K \times B)}$ where $h_1, h_2 \neq K$.

The result of these blocks depend on the result in Phase 2 and itself.

For instance, in the 1st iteration, the result of $D_{(2,3)}^{(2)}$ depends on $D_{(2,1)}^{(2)}$ and $D_{(1,3)}^{(2)}$:

$$d^{(1)}(3,5) = \min \left(d^{(0)}(3,5), d^{(2)}(3,1) + d^{(2)}(1,5) \right)$$

$$d^{(1)}(3,6) = \min \left(d^{(0)}(3,6), d^{(2)}(3,1) + d^{(2)}(1,6) \right)$$

$$d^{(1)}(4,5) = \min \left(d^{(0)}(4,5), d^{(2)}(4,1) + d^{(2)}(1,5) \right)$$

$$d^{(1)}(4,6) = \min \left(d^{(0)}(4,6), d^{(2)}(4,1) + d^{(2)}(1,6) \right)$$

$$d^{(2)}(3,5) = \min \left(d^{(1)}(3,5), d^{(2)}(3,2) + d^{(2)}(2,5) \right)$$

$$d^{(2)}(3,6) = \min \left(d^{(1)}(3,6), d^{(2)}(3,2) + d^{(2)}(2,6) \right)$$

$$d^{(2)}(4,5) = \min \left(d^{(1)}(4,5), d^{(2)}(4,2) + d^{(2)}(2,5) \right)$$

$$d^{(2)}(4,6) = \min \left(d^{(1)}(4,6), d^{(2)}(4,2) + d^{(2)}(2,6) \right)$$

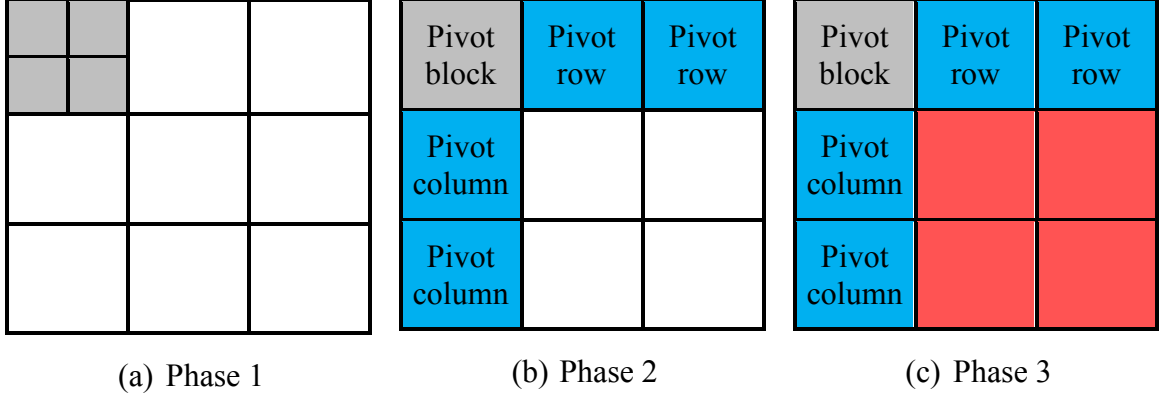


Figure 2: The 3 phases of blocked FW algorithm in the 1st iteration

The computations of $D_{(1,3)}^{(2)}$, $D_{(2,3)}^{(2)}$ and its dependencies are illustrated in Figure 3.

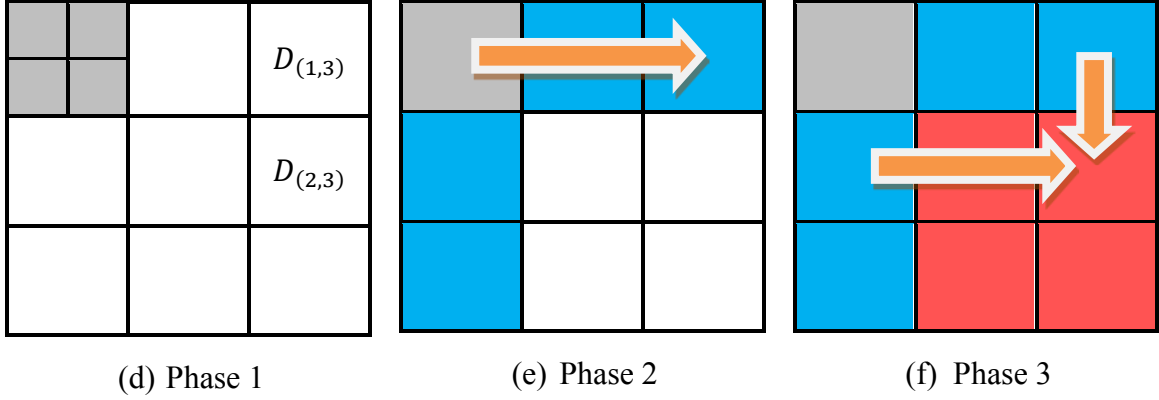


Figure 3: Dependencies of $D_{(1,3)}^{(2)}$, $D_{(2,3)}^{(2)}$ in the 1st iteration

In this particular example where $N = 6$ and $B = 2$, we will require $\lceil N/B \rceil = 3$ rounds.

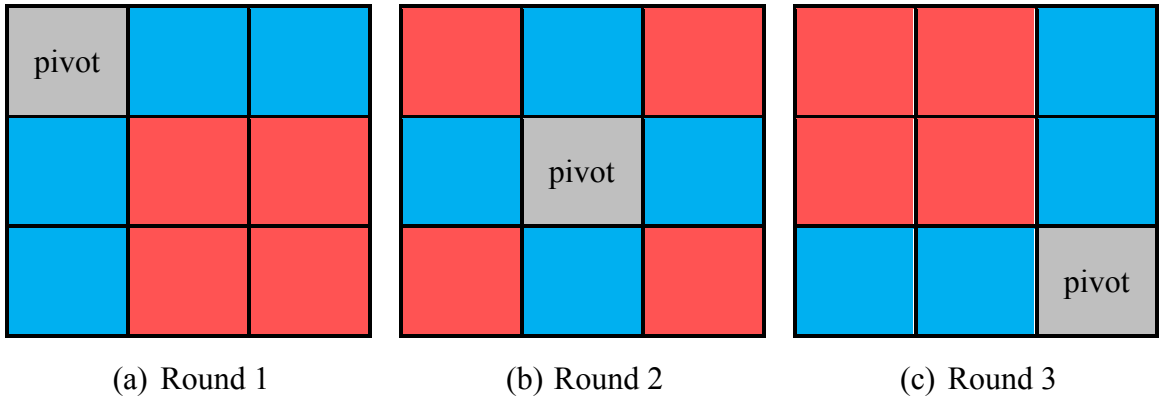


Figure 4: Blocked FW algorithm in each iteration

3 INPUT / OUTPUT FORMAT

1. Your program is required to read an input file, and generate output in another file.
2. Your program accepts 2 input parameters. They are:

- i 、 (String) the input file name
- ii 、 (String) the output file name
- iii 、 (Integer) the blocking factor

Make sure users can assign test cases through command line. For instance:

```
$ ./executable in_file out_file 32
```

TAs will judge your program as follows:

```
$ diff -b out_file answer
```

3. The 1st line of an input test case consists of 2 integers N ($1 \leq N \leq 10000$) and M ($0 \leq M \leq 10^9$) separated by a single space, which represents number of vertices and number of edge weight assignments respectively.

Each of the following M lines consists of 3 integers i, j and W ($i \neq j$), separated by a single space between any two numbers.

- i represents the index of the source vertex ($1 \leq i \leq N$)
- j represents the index of the destination vertex ($1 \leq j \leq N$)
- W represents the distance (weight of edge) from vertex i to vertex j ($0 \leq W \leq 100$)

Edges which are not listed in the input file do not exist in the graph. That is, for all $i \neq j$, if edge(i, j) does not show up in the input at all, vertex i does not have an edge to vertex j . But since we are dealing with a **directed** graph, this does NOT imply that edge(j, i) is also non-existent.

Besides, if there are re-assignments of an edge, please follow the latest one.

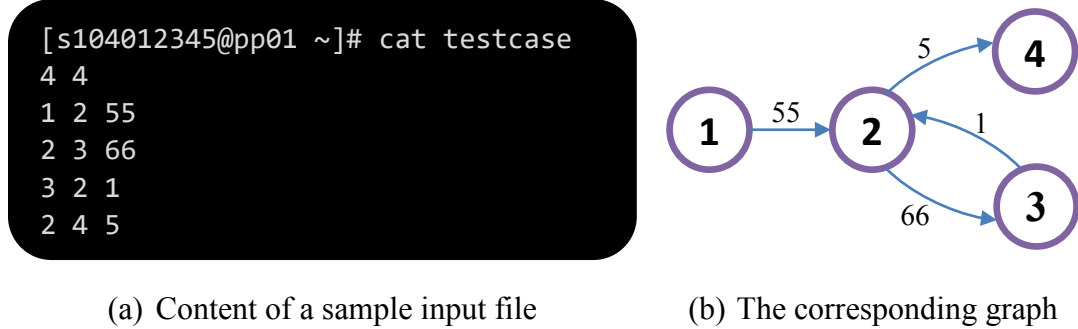


Figure 5: Sample Input

4. For output file, list the shortest-path distance of all vertex pairs.

Assume N represents total number of vertices, the output file should consists of N lines, each line consists of N numbers and separate them by a single space.

The number at the i^{th} line and the j^{th} column is the shortest-path distance from the i^{th} vertex to the j^{th} vertex if there is a path; otherwise, the corresponding output should be **INF**.

```
[s104012345@pp01 ~]# cat output
0 55 121 60
INF 0 66 5
INF 1 0 6
INF INF INF 0
```

Figure 6: Sample output

The sample test cases are provided in **/home/ipl2017/shared/hw4** on pp31.

4 WORKING ITEMS

You are required to implement **2** versions of blocked Floyd-Warshall algorithm under the given restrictions.

1. Single GPU

- Implement blocked APSP algorithm as described in Section 2.
- The main algorithm should be implemented in CUDA C/C++ kernel functions.
- Achieve better performance than sequential Floyd-Warshall implementation.

2. Multi GPUs implementation with OpenMP

- The restrictions of single-GPU version still hold.
- Able to utilize multiple GPUs available on single node.
- Achieve better performance than single GPU version.

3. Makefile

Please refer to the example in [/home/ipl2017/shared/hw4](#) on **apolloGPU**.
Don't modify execution file name(HW4_cuda.exe, HW4_openmp.exe) in sample Makefile.

4. README

You should specify your best block factor in README, TAs will use this configuration to test your performance.

Please refer to the example in [/home/ipl2017/shared/hw4](#) on **apolloGPU**.

5 OPTIMIZATION HINTS

- Shared memory
- Streaming
- Resolve bank conflicts
- Dynamic load-balancing (for openMP and MPI)

6 GRADING

1. Correctness (70%)

i 、 [50%] Single-GPU

ii 、 [20%] Multi-GPU implementation with OpenMP

2. Performance (20%)

- Performance is measured by the execution time of your program using 'time' Linux command.
- Points are giving according to the performance ranking of your program among all the students.

3. Demo (10%)

- Each student is given 10 minutes to explain your implementation followed by some questions from TA.
- No debugging or code modification is allowed during the demo.

- Points are given according to your understanding and explanation of your code, and your answers of the TA questions.

7 REMINDER

1. Please upload the following files to **HW_submission/HW4** directory on **apolloGPU** under your home directory before **8/3 8:00AM** (**The folder will be locked after deadline**)

i 、 **HW4_{student-ID}_cuda.cu**

ii 、 **HW4_{student-ID}_openmp.cu**

iii 、 **Makefile**

iv 、 **README**

Make sure your compile script can execute correctly and your code has no compile error before you upload your homework.

2. We provide sample code `seq_FW.cpp` and `block_FW.cpp` in **/home/ipl2017/shared/hw4** on **apolloGPU**.
3. Since we have limited resources for you guys to use, please start your work ASAP. Do not leave it until the last day!
4. **0 will be given to cheater** (even copying code from the Internet), but discussion on code is encouraged.
5. Asking questions through iLMS or by email are welcomed!